



UiO : **Universitetet i Oslo**

Logistic Regression and Neural Network

Logistic Regression and Neural Network

Author:

Ana Costa

authors@email.com

Supervisors:

UiO

November 13, 2018

Contents

1	General structure	3
1.1	Part a) Producing the data for the one-dimensional Ising model	3
1.2	Part b) Estimating the coupling constant of the one-dimensional Ising model	6
1.3	Understanding the results	7
1.4	Part c) Determine the phase of the two-dimensional Ising model.	13
1.5	Part d) Regression analysis of the one-dimensional Ising model using neural networks.	18
1.6	Part e) Classifying the Ising model phase using neural networks.	18
1.7	Part f) Critical evaluation of the various algorithms.	21
2	References	22

List of Figures

List of Tables

List of Codes

1.1	Code example	6
1.2	Performance	7
1.3	Performance MSE	8
1.4	Performance bias/variance	9
1.5	Performance OLS bias/variance	10
1.6	Performance Ridge bias/variance	11
1.7	Performance Lasso bias/variance	12
1.8	Plotting some states	15
1.9	Accuracy	16

<https://github.com/anacost/project2-FYS-STK4155>

1 General structure

- Project title
- Name, email, course title, date, group assistant
- Abstract (1/2 page max)
- Introduction (1 page)
- Method
 - Packages used
 - Datasets (models and observations)
 - Analysis method
 - ...
- Results
- Discussion and outlook (1 page)
- Conclusions (1/2 page)
- References
- Acknowledgments

```
1 import tensorflow
2 import sklearn
3 import pandas
4 import matplotlib.pyplot as plt
5 from mpl_toolkits.axes_grid1 import make_axes_locatable
6 import seaborn
7 %matplotlib inline
8 import numpy as np
```

1.1 Part a) Producing the data for the one-dimensional Ising model

```
1 import scipy.sparse as sp
2 np.random.seed(12)
3 import warnings
4 #Comment this to turn on warnings
5 warnings.filterwarnings('ignore')
6
7 ### define Ising model aprams
8 # system size
9 L=40
10
11 # create 10000 random Ising states
12 states=np.random.choice([-1, 1], size=(10000,L))
13
14 def ising_energies(states,L):
15     """
16     This function calculates the energies of the states in the nn
17     ↪ Ising Hamiltonian
18     """
19     J=np.zeros((L,L),)
20     for i in range(L):
21         J[i,(i+1)%L]=-1.0
22     # compute energies
23     E = np.einsum('...i,ij,...j->...',states,J,states)
```

```
24     return E
25 # calculate Ising energies
26 energies=ising_energies(states,L)
```

```
1 # reshape Ising states into RL samples: S_iS_j --> X_p
2 states=np.einsum('...i,...j->...ij', states, states)
3 shape=states.shape
4 states=states.reshape((shape[0],shape[1]*shape[2]))
5 # build final data set
6 Data=[states,energies]
```

```
1 # define number of samples
2 n_samples=400
3 # define train and test data sets
4 X_train=Data[0][:n_samples]
5 Y_train=Data[1][:n_samples] #+ np.random.normal(0,4.0,size=X_train.
    ↪ shape[0])
6 X_test=Data[0][n_samples:3*n_samples//2]
7 Y_test=Data[1][n_samples:3*n_samples//2] #+ np.random.normal(0,4.0,
    ↪ size=X_test.shape[0])
```


1.2 Part b) Estimating the coupling constant of the one-dimensional Ising model

Code 1.1: Code example

```
1  from sklearn import linear_model
2  # define error lists
3  train_errors_leastsq = []
4  test_errors_leastsq = []
5  train_MSE_leastsq = []
6  test_MSE_leastsq = []
7  train_bias_leastsq = []
8  test_bias_leastsq = []
9  train_var_leastsq = []
10 test_var_leastsq = []
11
12 train_errors_ridge = []
13 test_errors_ridge = []
14 train_MSE_ridge = []
15 test_MSE_ridge = []
16 train_bias_ridge = []
17 test_bias_ridge = []
18 train_var_ridge = []
19 test_var_ridge = []
20
21 train_errors_lasso = []
22 test_errors_lasso = []
23 train_MSE_lasso = []
24 test_MSE_lasso = []
25 train_bias_lasso = []
26 test_bias_lasso = []
27 train_var_lasso = []
28 test_var_lasso = []
29
30 # set regularisation strength values
31 lmbdas = np.logspace(-4, 5, 10)
32
33 #Initialize coefficients for OLS, ridge regression and Lasso
34 coefs_leastsq = []
35 coefs_ridge = []
36 coefs_lasso = []
37 # set up Lasso Regression model
38 lasso = linear_model.Lasso()
39
40 for _, lambda in enumerate(lmbdas):
41     ### ordinary least squares
42     xb = np.c_[np.ones((X_train.shape[0], 1)), X_train]
43     #fit model/singularity :
44     beta_ols = np.linalg.pinv(xb.T @ xb) @ xb.T @ Y_train
45     coefs_leastsq.append(beta_ols) # store weights
46
47     # use the coefficient of determination R^2 as the performance
48     # of prediction.
49     fitted_train = xb @ beta_ols
50     xb_test = np.c_[np.ones((X_test.shape[0], 1)), X_test]
51
52     fitted_test = xb_test @ beta_ols
53     R2_train = 1 - np.sum((fitted_train - Y_train)**2) / np.sum((Y_train - np.mean(Y_train))**2)
54     R2_test = 1 - np.sum((fitted_test - Y_test)**2) / np.sum((Y_test - np.mean(Y_test))**2)
55     MSE_train = np.sum((fitted_train - Y_train)**2) / len(Y_train)
56     MSE_test = np.sum((fitted_test - Y_test)**2) / len(Y_test)
57     var_train = np.sum((fitted_train - np.mean(fitted_train))**2) / len(Y_train)
```

1 General structure

6

Code 1.2: Performance

```
1 # Plot our performance on both the training and test data
2 plt.semilogx(lmbdas, train_errors_leastsq, 'b',label='Train (OLS)')
3 plt.semilogx(lmbdas, test_errors_leastsq,'--b',label='Test (OLS)')
4 plt.semilogx(lmbdas, train_errors_ridge,'r',label='Train (Ridge)',
5             ↳ linewidth=1)
6 plt.semilogx(lmbdas, test_errors_ridge,'--r',label='Test (Ridge)',
7             ↳ linewidth=1)
8
9 plt.semilogx(lmbdas, train_errors_lasso, 'g',label='Train (LASSO)')
10 plt.semilogx(lmbdas, test_errors_lasso, '--g',label='Test (LASSO)')
11
12 fig = plt.gcf()
13 fig.set_size_inches(10.0, 6.0)
14
15 #plt.vlines(alpha_optim, plt.ylim()[0], np.max(test_errors), color
16 ↳ ='k',
17
18 # linewidth=3, label='Optimum on test')
19 plt.legend(loc='lower left',fontsize=16)
20 plt.ylim([-0.01, 1.01])
21 plt.xlim([min(lmbdas), max(lmbdas)])
22 plt.xlabel(r'$\lambda$',fontsize=16)
23 plt.ylabel('Performance',fontsize=16)
24 plt.tick_params(labelsize=16)
25 plt.show()
```

1.3 Understanding the results

Let us make a few remarks: (i) the (inverse, see [Scikit documentation](#)) regularization parameter λ affects the Ridge and LASSO regressions at scales, separated by a few orders of magnitude. Notice that this is different for the data considered in Notebook 3 **Section VI: Linear Regression (Diabetes)**. Therefore, it is considered good practice to always check the performance for the given model and data with λ . (ii) at $\lambda \rightarrow 0$ and $\lambda \rightarrow \infty$, all three models overfit the data, as can be seen from the deviation of the test errors from unity (dashed lines), while the training curves stay at unity. (iii) While the OLS and Ridge regression test curves are monotonic, the LASSO test curve is not -- suggesting the optimal LASSO regularization parameter is $\lambda \approx 10^{-2}$. At this sweet spot, the Ising interaction weights \mathbf{J} contain only nearest-neighbor terms (as did the model the data was generated from).

Gauge degrees of freedom: recall that the uniform nearest-neighbor interactions strength $J_{j,k} = J$ which we used to generate the data was set to unity, $J = 1$. Moreover, $J_{j,k}$ was NOT defined to be symmetric (we only used the $J_{j,j+1}$ but never the $J_{j,j-1}$ elements). The colorbar on the matrix elements plot above suggest that the OLS and Ridge regression learn uniform symmetric weights $J = -0.5$. There is no mystery since this amounts to taking into account both the $J_{j,j+1}$ and the $J_{j,j-1}$ terms, and the weights are distributed symmetrically between them. LASSO, on the other hand, can break this symmetry (see matrix elements plots for $\lambda = 0.001$ and $\lambda = 0.01$). Thus, we see how different regularization schemes can lead to learning equivalent models but in different gauges. Any information we have about the symmetry of the unknown model that generated the data has to be reflected in the definition of the model and the regularization chosen.

Code 1.3: Performance MSE

```
1 # Plot our performance on both the training and test data
2 plt.semilogx(lmbdas, train_MSE_leastsq, 'b',label='Train (OLS)')
3 plt.semilogx(lmbdas, test_MSE_leastsq,'--b',label='Test (OLS)')
4 plt.semilogx(lmbdas, train_MSE_ridge,'r',label='Train (Ridge)',
5             ↪ linewidth=1)
6 plt.semilogx(lmbdas, test_MSE_ridge,'--r',label='Test (Ridge)',
7             ↪ linewidth=1)
8 plt.semilogx(lmbdas, train_MSE_lasso, 'g',label='Train (LASSO)')
9 plt.semilogx(lmbdas, test_MSE_lasso, '--g',label='Test (LASSO)')
10
11
12 fig = plt.gcf()
13 fig.set_size_inches(10.0, 6.0)
14
15 #plt.vlines(alpha_optim, plt.ylim()[0], np.max(test_errors), color
16             ↪ ='k',
17             # linewidth=3, label='Optimum on test')
18 plt.legend(loc='lower left',fontsize=16)
19 plt.ylim([-0.01, 1.01])
20 plt.xlim([min(lmbdas), max(lmbdas)])
21 plt.xlabel(r'$\lambda$',fontsize=16)
22 plt.ylabel('Performance-MSE',fontsize=16)
23 plt.tick_params(labelsize=16)
24 plt.show()
```


Code 1.4: Performance bias/variance

```
1 # Plot our bias-variance on both the training and test data
2 plt.semilogx(lmbdas, train_bias_leastsq, 'b',label='Bias-Train (OLS)'
   ↪ )
3 plt.semilogx(lmbdas, test_bias_leastsq,'--b',label='Bias-Test (OLS)'
   ↪ )
4 plt.semilogx(lmbdas, train_bias_ridge,'r',label='Bias-Train (Ridge)'
   ↪ ,linewidth=1)
5 plt.semilogx(lmbdas, test_bias_ridge,'--r',label='Bias-Test (Ridge)'
   ↪ ,linewidth=1)
6 plt.semilogx(lmbdas, train_bias_lasso, 'g',label='Bias-Train (LASSO)'
   ↪ )
7 plt.semilogx(lmbdas, test_bias_lasso, '--g',label='Bias-Test (LASSO)'
   ↪ )
8
9 plt.semilogx(lmbdas, train_var_leastsq, ':b',label='Variance-Train (
   ↪ OLS)')
10 plt.semilogx(lmbdas, test_var_leastsq, '.b',label='Variance-Test (OLS'
   ↪ )')
11 plt.semilogx(lmbdas, train_var_ridge, ':r',label='Variance-Train (
   ↪ Ridge)',linewidth=1)
12 plt.semilogx(lmbdas, test_var_ridge, '.r',label='Variance-Test (Ridge'
   ↪ )',linewidth=1)
13 plt.semilogx(lmbdas, train_var_lasso, ':g',label='Variance-Train (
   ↪ LASSO)')
14 plt.semilogx(lmbdas, test_var_lasso, '.g',label='Variance-Test (
   ↪ LASSO)')
15
16 fig = plt.gcf()
17 fig.set_size_inches(10.0, 6.0)
18
19 #plt.vlines(alpha_optim, plt.ylim()[0], np.max(test_errors), color
   ↪ ='k',
20 #           linewidth=3, label='Optimum on test')
21 plt.legend(loc='lower left',fontsize=16)
22 #plt.ylim([-0.01, 1.01])
23 plt.xlim([min(lmbdas), max(lmbdas)])
24 plt.xlabel(r'$\lambda$',fontsize=16)
25 plt.ylabel('Bias-Variance',fontsize=16)
26 plt.tick_params(labelsize=16)
27 plt.show()
```

Code 1.5: Performance OLS bias/variance

```
1 # Plot our bias-variance on both the training and test data
2 plt.semilogx(lmbdas, train_bias_leastsq, 'b',label='Bias-Train (OLS)
   ↪ ')
3 plt.semilogx(lmbdas, test_bias_leastsq,'--b',label='Bias-Test (OLS)'
   ↪ )
4 #plt.semilogx(lmbdas, train_bias_ridge,'r',label='Bias-Train (Ridge)
   ↪ ',linewidth=1)
5 #plt.semilogx(lmbdas, test_bias_ridge,'--r',label='Bias-Test (Ridge)
   ↪ ',linewidth=1)
6 #plt.semilogx(lmbdas, train_bias_lasso, 'g',label='Bias-Train (LASSO
   ↪ )')
7 #plt.semilogx(lmbdas, test_bias_lasso, '--g',label='Bias-Test (LASSO
   ↪ )')
8
9 plt.semilogx(lmbdas, train_var_leastsq, ':b',label='Variance-Train (
   ↪ OLS)')
10 plt.semilogx(lmbdas, test_var_leastsq, '.b',label='Variance-Test (OLS
   ↪ )')
11 #plt.semilogx(lmbdas, train_var_ridge, ':r',label='Variance-Train (
   ↪ Ridge)',linewidth=1)
12 #plt.semilogx(lmbdas, test_var_ridge, '.r',label='Variance-Test (
   ↪ Ridge)',linewidth=1)
13 #plt.semilogx(lmbdas, train_var_lasso, ':g',label='Variance-Train (
   ↪ LASSO)')
14 #plt.semilogx(lmbdas, test_var_lasso, '.g',label='Variance-Test (
   ↪ LASSO)')
15
16 fig = plt.gcf()
17 fig.set_size_inches(10.0, 6.0)
18
19 #plt.vlines(alpha_optim, plt.ylim()[0], np.max(test_errors), color
   ↪ ='k',
20 #           linewidth=3, label='Optimum on test')
21 plt.legend(loc='lower left',fontsize=16)
22 #plt.ylim([-0.01, 1.01])
23 plt.xlim([min(lmbdas), max(lmbdas)])
24 plt.xlabel(r'$\lambda$',fontsize=16)
25 plt.ylabel('Bias-Variance',fontsize=16)
26 plt.tick_params(labelsize=16)
27 plt.show()
```

Code 1.6: Performance Ridge bias/variance

```
1 # Plot our bias-variance on both the training and test data
2 #plt.semilogx(lmbdas, train_bias_leastsq, 'b',label='Bias-Train (OLS
   ↪ )')
3 #plt.semilogx(lmbdas, test_bias_leastsq,'--b',label='Bias-Test (OLS
   ↪ )')
4 plt.semilogx(lmbdas, train_bias_ridge,'r',label='Bias-Train (Ridge)'
   ↪ ,linewidth=1)
5 plt.semilogx(lmbdas, test_bias_ridge,'--r',label='Bias-Test (Ridge)'
   ↪ ,linewidth=1)
6 #plt.semilogx(lmbdas, train_bias_lasso, 'g',label='Bias-Train (LASSO
   ↪ )')
7 #plt.semilogx(lmbdas, test_bias_lasso, '--g',label='Bias-Test (LASSO
   ↪ )')
8
9 #plt.semilogx(lmbdas, train_var_leastsq, ':b',label='Variance-Train
   ↪ (OLS)')
10 #plt.semilogx(lmbdas, test_var_leastsq, '.b',label='Variance-Test (
   ↪ OLS)')
11 plt.semilogx(lmbdas, train_var_ridge, ':r',label='Variance-Train (
   ↪ Ridge)',linewidth=1)
12 plt.semilogx(lmbdas, test_var_ridge, '.r',label='Variance-Test (Ridge
   ↪ )',linewidth=1)
13 #plt.semilogx(lmbdas, train_var_lasso, ':g',label='Variance-Train (
   ↪ LASSO)')
14 #plt.semilogx(lmbdas, test_var_lasso, '.g',label='Variance-Test (
   ↪ LASSO)')
15
16 fig = plt.gcf()
17 fig.set_size_inches(10.0, 6.0)
18
19 #plt.vlines(alpha_optim, plt.ylim()[0], np.max(test_errors), color
   ↪ ='k',
20 #           linewidth=3, label='Optimum on test')
21 plt.legend(loc='lower left',fontsize=16)
22 #plt.ylim([-0.01, 1.01])
23 plt.xlim([min(lmbdas), max(lmbdas)])
24 plt.xlabel(r'$\lambda$',fontsize=16)
25 plt.ylabel('Bias-Variance',fontsize=16)
26 plt.tick_params(labelsize=16)
27 plt.show()
```

Code 1.7: Performance Lasso bias/variance

```
1 # Plot our bias-variance on both the training and test data
2 #plt.semilogx(lmbdas, train_bias_leastsq, 'b',label='Bias-Train (OLS
   ↪ ')')
3 #plt.semilogx(lmbdas, test_bias_leastsq,'--b',label='Bias-Test (OLS
   ↪ ')')
4 #plt.semilogx(lmbdas, train_bias_ridge,'r',label='Bias-Train (Ridge)
   ↪ ',linewidth=1)
5 #plt.semilogx(lmbdas, test_bias_ridge,'--r',label='Bias-Test (Ridge)
   ↪ ',linewidth=1)
6 plt.semilogx(lmbdas, train_bias_lasso, 'g',label='Bias-Train (LASSO)
   ↪ ')
7 plt.semilogx(lmbdas, test_bias_lasso, '--g',label='Bias-Test (LASSO)
   ↪ ')
8
9 #plt.semilogx(lmbdas, train_var_leastsq, ':b',label='Variance-Train
   ↪ (OLS)')
10 #plt.semilogx(lmbdas, test_var_leastsq, '.b',label='Variance-Test (
   ↪ OLS)')
11 #plt.semilogx(lmbdas, train_var_ridge, ':r',label='Variance-Train (
   ↪ Ridge)',linewidth=1)
12 #plt.semilogx(lmbdas, test_var_ridge, '.r',label='Variance-Test (
   ↪ Ridge)',linewidth=1)
13 plt.semilogx(lmbdas, train_var_lasso, ':g',label='Variance-Train (
   ↪ LASSO)')
14 plt.semilogx(lmbdas, test_var_lasso, '.g',label='Variance-Test (
   ↪ LASSO)')
15
16 fig = plt.gcf()
17 fig.set_size_inches(10.0, 6.0)
18
19 #plt.vlines(alpha_optim, plt.ylim()[0], np.max(test_errors), color
   ↪ ='k',
20 #           linewidth=3, label='Optimum on test')
21 plt.legend(loc='lower left',fontsize=16)
22 #plt.ylim([-0.01, 1.01])
23 plt.xlim([min(lmbdas), max(lmbdas)])
24 plt.xlabel(r'$\lambda$',fontsize=16)
25 plt.ylabel('Bias-Variance',fontsize=16)
26 plt.tick_params(labelsize=16)
27 plt.show()

1 #bootstrap:
```

1.4 Part c) Determine the phase of the two-dimensional Ising model.

```
1 np.random.seed(1) # shuffle random seed generator
2
3 # Ising model parameters
4 L=40 # linear system size
5 J=-1.0 # Ising interaction
6 T=np.linspace(0.25,4.0,16) # set of temperatures
7 T_c=2.26 # Onsager critical temperature in the TD limit
```

```

1 ##### prepare training and test data sets
2 import pickle,os
3 from sklearn.model_selection import train_test_split
4
5 ##### define ML parameters
6 num_classes=2
7 train_to_test_ratio=0.5 # training samples
8
9 # path to data directory
10 path_to_data=os.path.expanduser('.')+'/data/'
11
12 # load data
13 file_name = "Ising2DFM_reSample_L40_T=All.pkl" # this file contains
    ↳ 16*10000 samples taken in T=np.arange(0.25,4.0001,0.25)
14 data = pickle.load(open(path_to_data+file_name,'rb')) # pickle
    ↳ reads the file and returns the Python object (1D array,
    ↳ compressed bits)
15 data = np.unpackbits(data).reshape(-1, 1600) # Decompress array and
    ↳ reshape for convenience
16 data=data.astype('int')
17 data[np.where(data==0)]=-1 # map 0 state to -1 (Ising variable can
    ↳ take values +/-1)
18
19 file_name = "Ising2DFM_reSample_L40_T=All_labels.pkl" # this file
    ↳ contains 16*10000 samples taken in T=np.arange
    ↳ (0.25,4.0001,0.25)
20 labels = pickle.load(open(path_to_data+file_name,'rb')) # pickle
    ↳ reads the file and returns the Python object (here just a 1D
    ↳ array with the binary labels)
21
22 # divide data into ordered, critical and disordered
23 X_ordered=data[:70000,:]
24 Y_ordered=labels[:70000]
25
26 X_critical=data[70000:100000,:]
27 Y_critical=labels[70000:100000]
28
29 X_disordered=data[100000:,:]
30 Y_disordered=labels[100000:]
31
32 #X_ordered[np.where(X_ordered==0)]=-1 # map 0 state to -1 (Ising
    ↳ variable can take values +/-1)
33 #X_critical[np.where(X_critical==0)]=-1 # map 0 state to -1 (Ising
    ↳ variable can take values +/-1)
34 #X_disordered[np.where(X_disordered==0)]=-1 # map 0 state to -1 (
    ↳ Ising variable can take values +/-1)
35 del data,labels
36
37 # define training and test data sets
38 X=np.concatenate((X_ordered,X_disordered))
39 Y=np.concatenate((Y_ordered,Y_disordered))
40
41 # pick random data points from ordered and disordered states
42 # to create the training and test sets
43 X_train,X_test,Y_train,Y_test=train_test_split(X,Y,train_size=
    ↳ train_to_test_ratio)
44

```

14 # full data set

```

46 X=np.concatenate((X_critical,X))
47 Y=np.concatenate((Y_critical,Y))
48
49 print('X_train shape:', X_train.shape)
50 print('Y_train shape:', Y_train.shape)
51 print()

```

```
X_train shape: (65000, 1600)
Y_train shape: (65000,)
```

```
65000 train samples
30000 critical samples
65000 test samples
```

Code 1.8: Plotting some states

```
1 ##### plot a few Ising states
2
3 # set colourbar map
4 cmap_args=dict(cmap='plasma_r')
5
6 # plot states
7 fig, axarr = plt.subplots(nrows=1, ncols=3)
8
9 axarr[0].imshow(X_ordered[20001].reshape(L,L),**cmap_args)
10 axarr[0].set_title('$\\mathrm{ordered}\\ phase$', fontsize=16)
11 axarr[0].tick_params(labelsize=16)
12
13 axarr[1].imshow(X_critical[10001].reshape(L,L),**cmap_args)
14 axarr[1].set_title('$\\mathrm{critical}\\ region$', fontsize=16)
15 axarr[1].tick_params(labelsize=16)
16
17 im=axarr[2].imshow(X_disordered[50001].reshape(L,L),**cmap_args)
18 axarr[2].set_title('$\\mathrm{disordered}\\ phase$', fontsize=16)
19 axarr[2].tick_params(labelsize=16)
20
21 fig.subplots_adjust(right=2.0)
22
23 plt.show()
```

Code 1.9: Accuracy

```
1 ##### apply logistic regression
2 from sklearn import linear_model
3 from sklearn.neural_network import MLPClassifier
4
5 # define regularisation parameter
6 lmbdas=np.logspace(-5,5,11)
7
8 # preallocate data
9 train_accuracy=np.zeros(lmbdas.shape,np.float64)
10 test_accuracy=np.zeros(lmbdas.shape,np.float64)
11 critical_accuracy=np.zeros(lmbdas.shape,np.float64)
12
13 train_accuracy_SGD=np.zeros(lmbdas.shape,np.float64)
14 test_accuracy_SGD=np.zeros(lmbdas.shape,np.float64)
15 critical_accuracy_SGD=np.zeros(lmbdas.shape,np.float64)
16
17 # loop over regularisation strength
18 for i,lambda in enumerate(lmbdas):
19
20     # define logistic regressor
21     logreg=linear_model.LogisticRegression(C=1.0/lambda,random_state
    ↪ =1,verbose=0,max_iter=1E3,tol=1E-5)
22
23     # fit training data
24     logreg.fit(X_train, Y_train)
25
26     # check accuracy
27     train_accuracy[i]=logreg.score(X_train,Y_train)
28     test_accuracy[i]=logreg.score(X_test,Y_test)
29     critical_accuracy[i]=logreg.score(X_critical,Y_critical)
30
31     print('accuracy: train, test, critical')
32     print('liblin: %0.4f, %0.4f, %0.4f' %(train_accuracy[i],
    ↪ test_accuracy[i],critical_accuracy[i]))
33
34     # define SGD-based logistic regression
35     logreg_SGD = linear_model.SGDClassifier(loss='log', penalty='l2
    ↪ ', alpha=lambda, max_iter=100,
36
37                                     shuffle=True,
38     ↪ random_state=1, learning_rate='optimal')
39
40     # fit training data
41     logreg_SGD.fit(X_train,Y_train)
42
43     # check accuracy
44     train_accuracy_SGD[i]=logreg_SGD.score(X_train,Y_train)
45     test_accuracy_SGD[i]=logreg_SGD.score(X_test,Y_test)
46     critical_accuracy_SGD[i]=logreg_SGD.score(X_critical,Y_critical
    ↪ )
47
48     print('SGD: %0.4f, %0.4f, %0.4f' %(train_accuracy_SGD[i],
    ↪ test_accuracy_SGD[i],critical_accuracy_SGD[i]))
49
50     print('finished computing %i/11 iterations' %(i+1))
51
52 # plot accuracy against regularisation strength
53 plt.semilogx(lmbdas,train_accuracy,'*-b',label='liblinear train')
54 plt.semilogx(lmbdas,test_accuracy,'*-r',label='liblinear test')
55 plt.semilogx(lmbdas,critical_accuracy,'*-g',label='liblinear
    ↪ critical')
56
57 plt.semilogx(lmbdas,train_accuracy_SGD,'*--b',label='SGD train')
```



```

1 ##### apply logistic regression
2 import logisRegresANA
3 import importlib
4 importlib.reload(logisRegresANA)
5 lr = 0.5
6 epochs = 300
7 weights= logisRegresANA.logistic_reg(X_train, Y_train, epochs, lr)

```

```

1 ##### apply logistic regression
2 import logisRegresANA
3 import importlib
4 importlib.reload(logisRegresANA)
5 weights=logisRegresANA.stocGradAscentA(X_train,Y_train)

```

```

1 ##### apply logistic regression
2 import logisRegresANA
3 import importlib
4 importlib.reload(logisRegresANA)
5 weights = logisRegresANA.steepest_descent_auto(X_train, Y_train,
    ↪ alpha =0.001)
6 error_train = logisRegresANA.simpctest(weights, X_train, Y_train)
7 error_test = logisRegresANA.simpctest(weights, X_test, Y_test)

```

```

1 ##### apply logistic regression
2 import logisRegresANA
3 import importlib
4 importlib.reload(logisRegresANA)
5 weights = logisRegresANA.logistic_reg(X_train, Y_train, epochs=
    ↪ 100, lr=0.001)
6 error_train = logisRegresANA.simpctest(weights, X_train, Y_train)
7 error_test = logisRegresANA.simpctest(weights, X_test, Y_test)

```

```

1 ##### apply logistic regression
2 import logisRegresANA
3 import importlib
4 importlib.reload(logisRegresANA)
5 weights = logisRegresANA.gradDscent(X_train, Y_train, alpha= 0.01)

```

```

1 import importlib
2 import logisRegresANA
3 importlib.reload(logisRegresANA)
4
5 weights2 = logisRegresANA.sgd(X_train, Y_train)
6 print(weights2)
7 weights2 = weights2.flatten()
8 error_train = logisRegresANA.simpctest(weights2, X_train, Y_train)
9 error_test = logisRegresANA.simpctest(weights2, X_test, Y_test)

```

1.5 Part d) Regression analysis of the one-dimensional Ising model using neural networks.

1.6 Part e) Classifying the Ising model phase using neural networks.

```
1 import importlib
2 import logisRegresANA
3 importlib.reload(logisRegresANA)
4
5 in_layer = X_train.shape[1] #number of neurons in the input layer
6
7 if (len(Y_train.shape)==1):
8     out_layer = 1 #number of neurons in the output layer
9 else: out_layer = Y_train.shape[1]
10 biasesnn, weightsnn= logisRegresANA.neuralnetwork([in_layer, 10,
11     ↪ out_layer], X_train, Y_train,
12                                     validation_x=
13     ↪ X_test, validation_y=Y_test,
14                                     verbose=True,
15                                     epochs= 30, mini_batch_size = 10, lr=
16     ↪ 0.5, C='ce')
```

```
mini_batch [[-1 -1 -1 ..., -1 -1  1]
[-1 -1 -1 ..., -1 -1  0]
[ 1  1  1 ...,  1  1  1]
...,
[ 1  1  1 ..., -1 -1  0]
[-1 -1 -1 ..., -1 -1  1]
[ 1  1 -1 ...,  1 -1  0]]
nabla_w from weights -shape (2,)
nabla_b from weights -shape (2,)
x [-1 -1 -1 ..., -1 -1 -1]
y 1
activations: [array([-1, -1, -1, ..., -1, -1, -1])]
f 0 w in loop [[ 0.50892274 -0.28898022 -1.20827711 ..., -0.456469 -
0.65965346
0.10241962]
[ 2.3897599  1.29156472  1.27318914 ..., -0.55426639 -0.94425406
-0.90373304]
[ 0.97948496 -0.55957186 -0.1711543 ...,  0.08315642 -0.76727449
-0.1474379 ]
...,
[-0.75724226 -0.53077548 -1.08823808 ..., -0.28865485  0.48977372
1.44011092]
[-0.01027559  0.83852329 -0.11277865 ..., -0.85940721  0.45309176
-0.42789696]
```

```

[ 0.26413263  1.05422235 -0.36335701 ..., -1.30188984 -0.40503143
 -0.68552322]]
activation : [-1 -1 -1 ..., -1 -1 -1]
w_a = numpy.dot(w , activation): [-13.49316819 -39.12589526 -
10.07647132  37.48403681 -26.31504817
 15.09851548 -13.01447302  18.43597532 -29.68447902 -0.04488525]
b : [-0.9073413  -1.36768501 -0.35959421 -0.89205373 -1.35507506 -
0.08539942
 -1.56692176  0.74175751  0.30745289 -0.95706228]
z = w_a + b : [-14.40050949 -40.49358026 -10.43606554  36.59198309 -
27.67012324
 15.01311605 -14.58139478  19.17773283 -29.37702613 -1.00194753]
f 1 w in loop [[ 0.74683728 -1.36658155  1.73000956 -0.99789933 -
0.8230627  -0.840913
 -0.48278465  0.37804188  0.20443838  0.55033394]]
activation : [ 5.57106144e-07  2.59335242e-18  2.93536138e-05
1.00000000e+00
 9.61652500e-13  9.99999698e-01  4.64922440e-07  9.99999995e-01
 1.74469556e-13  2.68558686e-01]
w_a = numpy.dot(w , activation): [-1.31292226]
b : [ 0.6802191]
z = w_a + b : [-0.63270316]
ACTIVATIONS : [-1 -1 -1 ..., -1 -1 -1]
[ 5.57106144e-07  2.59335242e-18  2.93536138e-05  1.00000000e+00
 9.61652500e-13  9.99999698e-01  4.64922440e-07  9.99999995e-01
 1.74469556e-13  2.68558686e-01]
[ 0.34689786]
in last layer, y 1 activations[-1 ] [ 0.34689786]
zs stored : [array([-14.40050949, -40.49358026, -10.43606554,
36.59198309,
 -27.67012324,  15.01311605, -14.58139478,  19.17773283,
 -29.37702613, -1.00194753]), array([-0.63270316])]
DELTA : [-0.65310214]
in backprop: numpy.array(activations[-2]).T [ 5.57106144e-07
2.59335242e-18  2.93536138e-05  1.00000000e+00
 9.61652500e-13  9.99999698e-01  4.64922440e-07  9.99999995e-01
 1.74469556e-13  2.68558686e-01]
in backprop: nabla_b_backprop[-1] [-0.65310214]
in backprop: nabla_w_backprop[-1] [ -3.63847217e-07 -1.69372402e-18 -
1.91709081e-05 -6.53102143e-01
 -6.28057308e-13 -6.53101946e-01 -3.03641842e-07 -6.53102140e-01
 -1.13946441e-13 -1.75396253e-01]
ENTRA NESTE LOOP, k = 1
sp [ 5.57105834e-07  2.59335242e-18  2.93527522e-05  2.22044605e-
16
 9.61652500e-13  3.01916105e-07  4.64922224e-07  4.69047088e-09
 1.74469556e-13  1.96434918e-01]
numpy.array(weights[-1- (k-1)]).T [[ 0.74683728]
[-1.36658155]
[ 1.73000956]
[-0.99789933]
[-0.8230627 ]
[-0.840913 ]
[-0.48278465]
[ 0.37804188]

```

```
[ 0.20443838]
[ 0.55033394]]
delta -0.653102142795
delta in loop k [ -2.71734513e-07    2.31461199e-18   -3.31648807e-05
1.44713173e-16
    5.16930545e-13    1.65812954e-07    1.46593551e-07   -1.15807709e-09
-2.32950262e-14   -7.06034783e-02]
testyy [-1 -1 -1 ..., -1 -1 -1]
```

ValueError
→ call last)

Traceback (most recent

```
<ipython-input-29-0374de8f71b8> in <module>
    11
→ validation_x=X_test, validation_y=Y_test,
    12                                     verbose=
→ True,
---> 13                                     epochs= 30, mini_batch_size =
→ 10, lr= 0.5, C='ce')
```

```
~/Documents/FYS-STK4155/Project2/logisRegresANA.py in neuralnetwork
→ (sizes, X_train, Y_train, validation_x, validation_y, verbose,
→ epochs, mini_batch_size, lr, C)
    138     #print('initial weights ', weights)
    139     biases, weights = SGD(X_train, Y_train, epochs,
→ mini_batch_size, lr, C, sizes, num_layers,
--> 140     biases, weights, verbose, validation_x,
→ validation_y)
    141     return biases, weights
    142 def SGD(X_train, Y_train, epochs, mini_batch_size, lr, C,
→ sizes, num_layers,
```

```
~/Documents/FYS-STK4155/Project2/logisRegresANA.py in SGD(X_train,
→ Y_train, epochs, mini_batch_size, lr, C, sizes, num_layers,
→ biases, weights, verbose, validation_x, validation_y)
    152     #feed-forward (and back) all mini_batches
    153     for k, minib in enumerate(mini_batches):
--> 154         biases, weights = update_mini_batch(minib, lr,
→ C, sizes, num_layers, biases, weights)
    155
    156     if(verbose):
```

```
~/Documents/FYS-STK4155/Project2/logisRegresANA.py in
→ update_mini_batch(minibatch, lr, C, sizes, num_layers, biases,
→ weights)
    185         print('y ', y)
    186         #backpropagation for each observation in
→ mini_batch
--> 187         delta_nabla_b, delta_nabla_w = backprop(x, y, C
→ , sizes, num_layers, biases, weights)
    188         #print(delta_nabla_b.shape, 'delta_nabla_b ',
→ delta_nabla_b)
    189         print(delta_nabla_w.shape, 'delta_nabla_w ',
```

```
↪ delta_nabla_w)
```

```
~/Documents/FYS-STK4155/Project2/logisRegresANA.py in backprop(x, y
↪ , C, sizes, num_layers, biases, weights)
    261         testyy = numpy.array(activations[-1 - (k+1)]).T
    262         print('testyy', testyy)
--> 263         nabla_w_backprop[-1 -(k)] = numpy.multiply(
↪ delta , testyy)
    264
    265         return nabla_b_backprop, nabla_w_backprop
```

```
ValueError: operands could not be broadcast together with shapes
↪ (10,) (1600,)
```

1.7 Part f) Critical evaluation of the various algorithms.

2 References